

LoanTap Classification Project

1. Define Problem Statement and perform Exploratory Data Analysis

1.1. Definition of problem

The Goal of this project is to predict/make decision to give loan to customers to mitigate the risk of defaults

In [3]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import chi2_contingency, f_oneway, ttest_ind, spearmanr
from scipy.stats import chi2, shapiro, boxcox, ttest_rel
from scipy.stats import chisquare, kruskal
import statsmodels.api as sm
import scipy.stats as stats
```

In [4]:

```
df=pd.read_csv('https://d2beiqrkhq929f0.cloudfront.net/public_assets/assets/000/003/549/original/logistic_regression.csv?1651045921')
```

In [5]:

```
df.shape
```

Out[5]:

```
(396030, 27)
```

In [6]:

```
df.head()
```

Out[6]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	op
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	...	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...	

5 rows × 27 columns



In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                             396030 non-null float64
1   term                                  396030 non-null object
2   int_rate                             396030 non-null float64
3   installment                           396030 non-null float64
4   grade                                 396030 non-null object
5   sub_grade                             396030 non-null object
6   emp_title                             373103 non-null object
7   emp_length                           377729 non-null object
8   home_ownership                       396030 non-null object
9   annual_inc                           396030 non-null float64
10  verification_status                  396030 non-null object
11  issue_d                              396030 non-null object
12  loan_status                          396030 non-null object
13  purpose                              396030 non-null object
14  title                                394275 non-null object
15  dti                                   396030 non-null float64
16  earliest_cr_line                     396030 non-null object
17  open_acc                             396030 non-null float64
18  pub_rec                              396030 non-null float64
19  revol_bal                            396030 non-null float64
20  revol_util                           395754 non-null float64
21  total_acc                            396030 non-null float64
22  initial_list_status                  396030 non-null object
23  application_type                     396030 non-null object
24  mort_acc                             358235 non-null float64
25  pub_rec_bankruptcies                 395495 non-null float64
26  address                              396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

2.1.Duplicate Value Check

In [8]:

```
df.duplicated().sum()
```

Out[8]:
0

In [9]:

```
df.describe()
```

Out[9]:

	loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_b
count	396030.000000	396030.000000	396030.000000	3.960300e+05	396030.000000	396030.000000	396030.000000	3.960300e+05
mean	14113.888089	13.639400	431.849698	7.420318e+04	17.379514	11.311153	0.178191	1.584454e+04
std	8357.441341	4.472157	250.727790	6.163762e+04	18.019092	5.137649	0.530671	2.059184e+04
min	500.000000	5.320000	16.080000	0.000000e+00	0.000000	0.000000	0.000000	0.000000e+00
25%	8000.000000	10.490000	250.330000	4.500000e+04	11.280000	8.000000	0.000000	6.025000e+03
50%	12000.000000	13.330000	375.430000	6.400000e+04	16.910000	10.000000	0.000000	1.118100e+04
75%	20000.000000	16.490000	567.300000	9.000000e+04	22.980000	14.000000	0.000000	1.962000e+04
max	40000.000000	30.990000	1533.810000	8.706582e+06	9999.000000	90.000000	86.000000	1.743266e+06

1.2.Conversion of categorical attributes to 'category'

In [10]:

```
columns=df.select_dtypes(include=['object']).columns
for column in columns:
    print(f'column name :{column}')
    print(df[column].value_counts(ascending=True))
    print('.....')
    print('.....')
    print('.....')
```

column name :term

60 months 94025

36 months 302005

Name: term, dtype: int64

.....

.....

.....

column name :grade

G 3054

F 11772

E 31488

D 63524

A 64187

C 105987

B 116018

Name: grade, dtype: int64

.....

.....

.....

column name :sub_grade

G5 316

G4 374

G3 552

G2 754

G1 1058

F5 1397

F4 1787

F3 2286

F2 2766

F1 3536

E5 4572

E4 5361

E3 6207

E2 7431

E1 7917

A2 9567

D5 9700

A1 9729

A3 10576

D4 11657

D3 12223

D2 13951

A4 15789

D1 15993

C5 18244

A5 18526

B1 19182

C4 20280

C3 21221

B5 22085

B2 22495

C2 22580

C1 23662

B4 25601

B3 26655

Name: sub_grade, dtype: int64

```
.....
.....
.....
.....
column name :emp_title
analytic33          1
Odin fashion        1
Aztec Animal Clinic 1
MCCS                1
Hand Surgery & Rehabilitation Center 1
```

```
.....
Supervisor          1830
RN                  1846
Registered Nurse    1856
Manager             4250
Teacher             4389
Name: emp_title, Length: 173105, dtype: int64
```

```
.....
.....
.....
column name :emp_length
9 years            15314
8 years            19168
7 years            20819
6 years            20841
4 years            23952
1 year             25882
5 years            26495
3 years            31665
< 1 year           31725
2 years            35827
10+ years          126041
Name: emp_length, dtype: int64
```

```
.....
.....
column name :home_ownership
ANY                 3
NONE                31
OTHER               112
OWN                 37746
RENT                159790
MORTGAGE            198348
Name: home_ownership, dtype: int64
```

```
.....
column name :verification_status
Not Verified        125082
Source Verified     131385
Verified            139563
Name: verification_status, dtype: int64
```

```
.....
column name :issue_d
Jun-2007            1
Sep-2007            15
Nov-2007            22
Sep-2008            25
Jul-2007            26
.....
Nov-2013            10496
Dec-2013            10618
Jan-2015            11705
Jul-2014            12609
Oct-2014            14846
Name: issue_d, Length: 115, dtype: int64
```

.....
.....
.....
.....
column name :loan_status
Charged Off 77673
Fully Paid 318357
Name: loan_status, dtype: int64
.....
.....
.....

.....
column name :purpose
educational 257
renewable_energy 329
wedding 1812
house 2201
vacation 2452
moving 2854
medical 4196
car 4697
small_business 5701
major_purchase 8790
other 21185
home_improvement 24030
credit_card 83019
debt_consolidation 234507
Name: purpose, dtype: int64
.....
.....
.....

.....
column name :title
Appliances Falling Apart 1
Wedding1 1
Finance Relief 1
kevins loan 1
cridit payoff 1
.....
Debt Consolidation 11608
Other 12930
Home improvement 15264
Credit card refinancing 51487
Debt consolidation 152472
Name: title, Length: 48817, dtype: int64
.....
.....
.....

.....
column name :earliest_cr_line
Aug-1959 1
Nov-1957 1
Jul-1958 1
Apr-1960 1
Dec-1950 1
.....
Nov-2000 2736
Aug-2001 2884
Oct-2001 2896
Aug-2000 2935
Oct-2000 3017
Name: earliest_cr_line, Length: 684, dtype: int64
.....
.....
.....

.....
column name :initial_list_status
w 157964
f 238066
Name: initial_list_status, dtype: int64
.....
.....

```
.....
.....
column name :application_type
DIRECT_PAY      286
JOINT           425
INDIVIDUAL      395319
Name: application_type, dtype: int64
.....
.....
.....
column name :address
0174 Michelle Gateway\r\nMendozaberg, OK 22690      1
01246 Carrie Passage\r\nNew Kyle, ND 11650          1
16658 Perez Key\r\nCampbellside, AL 70466           1
110 Gomez Flat Apt. 561\r\nJosephchester, RI 70466  1
27963 Jessica Lodge Suite 200\r\nDunntown, CT 70466  1
..
USNS Johnson\r\nFPO AP 48052                        7
USS Smith\r\nFPO AP 70466                            8
USS Johnson\r\nFPO AE 48052                         8
USNS Johnson\r\nFPO AE 05113                       8
USCGC Smith\r\nFPO AE 70466                        8
Name: address, Length: 393700, dtype: int64
.....
.....
.....
```

Treatment of column 'term'

In [11]:

```
df['term']=df['term'].str.split(' ',expand=True)[1]
```

In [12]:

```
df['term'].value_counts()
```

Out[12]:

```
36    302005
60     94025
Name: term, dtype: int64
```

In [13]:

```
df.rename(columns = {'term':'term (months)'}, inplace = True)
```

In [14]:

```
df.head()
```

Out[14]:

	loan_amnt	term (months)	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	c
0	10000.0	36	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	
1	8000.0	36	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	
2	15600.0	36	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	
3	7200.0	36	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	...	

	loan_amnt	term (months)	int_rate	installment	grade	sub_grade	emp_title Destiny	emp_length	home_ownership	annual_inc	...	c
4	24375.0	60	17.27	609.33	C	C5	Management Inc.	9 years	MORTGAGE	55000.0	...	

5 rows x 27 columns



Treatment of column 'address'

In [15]:

```
df['address']=df['address'].str[-8:-6]
```

In [16]:

```
df['address'].value_counts()
```

Out[16]:

AP	14308
AE	14157
AA	13919
NJ	7091
WI	7081
LA	7068
NV	7038
AK	7034
MA	7022
VA	7022
VT	7005
NY	7004
MS	7003
TX	7000
SC	6973
ME	6972
AR	6969
OH	6969
GA	6967
ID	6958
IN	6958
KS	6945
WV	6944
RI	6940
MO	6939
IL	6934
WY	6933
NE	6927
HI	6927
IA	6926
FL	6921
AZ	6918
CO	6914
OK	6911
CT	6904
MN	6904
NC	6901
OR	6898
CA	6898
AL	6898
MD	6896
WA	6895
UT	6887
SD	6887
MT	6883
DE	6874
TN	6869
ND	6858
MI	6854
DC	6842
...	...

NM 6842
PA 6825
NH 6818
KY 6800
Name: address, dtype: int64

Treatment of column 'emp_length'

In [17]:

```
df['emp_length'].replace('< 1 year','0 year', inplace=True)
```

In [18]:

```
df['emp_length'].replace('10+ years','10 years', inplace=True)
```

In [19]:

```
df['emp_length']=df['emp_length'].str.split(' ',expand=True)[0]
```

In [20]:

```
df.rename(columns = {'emp_length':'emp_length (years)'}, inplace = True)
```

In [21]:

```
df.head()
```

Out[21]:

	loan_amnt	term (months)	int_rate	installment	grade	sub_grade	emp_title	emp_length (years)	home_ownership	annual_inc	...	c
0	10000.0	36	11.44	329.48	B	B4	Marketing	10	RENT	117000.0	...	
1	8000.0	36	11.99	265.68	B	B5	Credit analyst	4	MORTGAGE	65000.0	...	
2	15600.0	36	10.49	506.97	B	B3	Statistician	0	RENT	43057.0	...	
3	7200.0	36	6.49	220.65	A	A2	Client Advocate	6	RENT	54000.0	...	
4	24375.0	60	17.27	609.33	C	C5	Destiny Management Inc.	9	MORTGAGE	55000.0	...	

5 rows x 27 columns



2.2.Missing value treatment

In [22]:

```
df.isna().sum()
```

Out[22]:

loan_amnt 0
term (months) 0
int_rate 0
installment 0
grade 0
sub_grade 0
emp_title 22927
emp_length (years) 18301
home_ownership 0
annual_inc 0
verification_status 0
issue_d 0
loan status 0


```
purpose          0
title            1755
dti              0
earliest_cr_line 0
open_acc         0
pub_rec          0
revol_bal        0
revol_util       276
total_acc        0
initial_list_status 0
application_type 0
mort_acc         37795
pub_rec_bankruptcies 535
address          0
dtype: int64
```

Dropping rows with null value for 'pub_rec_bankruptcies', 'revol_util' & 'title' as they are less than 1% of the data

In [23]:

```
df = df.dropna(axis=0, subset=['pub_rec_bankruptcies', 'revol_util', 'title'])
```

In [24]:

```
df.isna().sum()
```

Out[24]:

```
loan_amnt          0
term (months)      0
int_rate           0
installment        0
grade              0
sub_grade          0
emp_title          22668
emp_length (years) 18076
home_ownership      0
annual_inc         0
verification_status 0
issue_d            0
loan_status        0
purpose            0
title              0
dti                0
earliest_cr_line   0
open_acc           0
pub_rec            0
revol_bal          0
revol_util         0
total_acc          0
initial_list_status 0
application_type    0
mort_acc           37195
pub_rec_bankruptcies 0
address            0
dtype: int64
```

Imputing 'mort_acc' by linear interpolation

In [25]:

```
df['mort_acc'] = df['mort_acc'].interpolate(method = 'linear',      limit_direction = 'forward', axis = 0)
```

In [26]:

```
df.isna().sum()
```

Out[26]:

```
loan_amnt      0
term (months)  0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22668
emp_length (years) 18076
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0
purpose        0
title          0
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     0
total_acc      0
initial_list_status 0
application_type 0
mort_acc       0
pub_rec_bankruptcies 0
address        0
dtype: int64
```

Imputing 'emp_length (years)' & 'emp_title' by forward fill by making dataframe in order by 'loan_amnt','annual_inc','installment','grade','sub_grade'

In [27]:

```
df = df.sort_values(by = ['loan_amnt','annual_inc', 'installment','grade','sub_grade'])
```

In [28]:

```
df.fillna(method="ffill", inplace=True)
```

In [29]:

```
df.isna().sum()
```

Out[29]:

```
loan_amnt      0
term (months)  0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      0
emp_length (years) 0
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0
purpose        0
title          0
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     0
total_acc      0
initial_liststatus 0
application_type 0
```

```
mort_acc      0
pub_rec_bankruptcies  0
address       0
dtype: int64
```

In [30]:

```
# shuffle the DataFrame rows
df = df.sample(frac = 1)
```

In [31]:

```
df.reset_index(drop=True)
```

Out[31]:

	loan_amnt	term (months)	int_rate	installment	grade	sub_grade	emp_title	emp_length (years)	home_ownership	annual_inc
0	5000.0	36	15.61	174.83	D	D1	Instructor	2	RENT	26000.0
1	12575.0	36	13.33	425.71	C	C3	Computer Network Engineer	10	RENT	127900.0
2	5125.0	36	16.78	182.16	C	C5	Administrative Assistant	10	MORTGAGE	44316.0
3	8000.0	36	14.09	273.78	B	B5	Indiana University Health	4	RENT	105000.0
4	9600.0	36	10.99	314.25	B	B3	General Manager	10	MORTGAGE	82940.0
...
393460	8825.0	36	15.61	308.57	D	D1	Psych Tech	10	MORTGAGE	26000.0
393461	1000.0	36	11.99	33.21	B	B3	Designer	4	OWN	30000.0
393462	6000.0	36	12.69	201.27	C	C2	Vice President, Human Resources	10	MORTGAGE	115000.0
393463	18000.0	36	12.99	606.41	B	B5	Sr. Account Mgr	10	MORTGAGE	170000.0
393464	18000.0	36	11.14	590.50	B	B2	Del Mar College	6	MORTGAGE	73000.0

393465 rows x 27 columns




1.3.Univariate Analysis


In [32]:

```
sns.boxplot(x=df['loan_amnt'])
```

Out[32]:

<AxesSubplot:xlabel='loan_amnt'>






Median Loan amount is 12000 and distribution is right skewed

In [33]:

```
sns.countplot(x=df['term (months)'])
```

Out[33]:

```
<AxesSubplot:xlabel='term (months)', ylabel='count'>
```




Around 75% of loans are having 3 Years of term

In [34]:

```
sns.boxplot(x=df['int_rate'])
```

Out[34]:

```
<AxesSubplot:xlabel='int_rate'>
```




Median interest rate is close to 13.5 and Distribution is right skewed


In [35]:

```
sns.boxplot(x=df['installment'])
```

Out[35]:

```
<AxesSubplot:xlabel='installment'>
```






Monthly median installment amount is close to 350 and Distribution is right skewed

In [36]:

```
sns.countplot(x=df['grade'],order=df['grade'].value_counts().iloc[:7].index)
```

Out[36]:

<AxesSubplot:xlabel='grade', ylabel='count'>



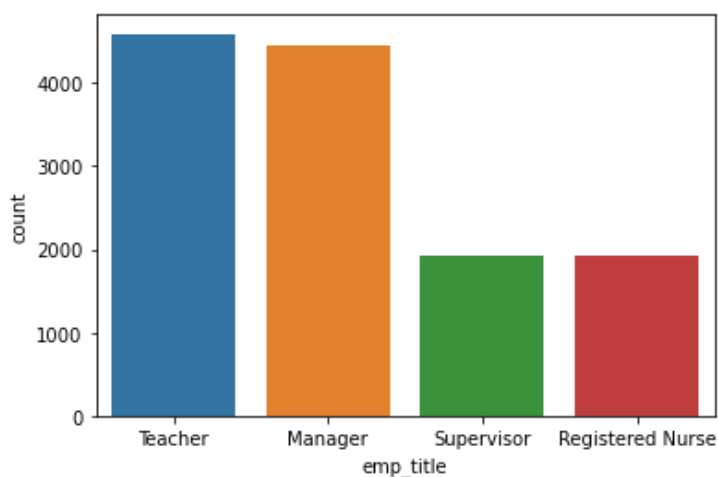
Maximum loan are graded as B category

In [37]:

```
sns.countplot(x=df['emp_title'],order=df['emp_title'].value_counts().iloc[:4].index)
```

Out[37]:

<AxesSubplot:xlabel='emp_title', ylabel='count'>




Teacher, Manager and Supervisor are amongst the top 3 employee title

In [38]:

```
sns.countplot(x=df['emp_length (years)'],order=df['emp_length (years)'].value_counts().iloc[:11].index)
```

Out[38]:

<AxesSubplot:xlabel='emp_length (years)', ylabel='count'>




Around 30% of Loans are given to employee with 10+years working lenght

In [39]:

```
sns.countplot(x=df['home_ownership'],order=df['home_ownership'].value_counts().iloc[:6].index)
```

Out[39]:

<AxesSubplot:xlabel='home_ownership', ylabel='count'>




Around 50% of home ownership is MORTGAGE

In [40]:

```
ax=sns.boxplot(x=df['annual_inc'])
```

```
plt.xlim([0, 400000])
```

```
plt.show()
```




Mean annual income is around 65000

In [41]:

```
sns.countplot(x=df['verification_status'])
```

Out[41]:

```
<AxesSubplot:xlabel='verification_status', ylabel='count'>
```




Not much difference in count of verification status

In [42]:

```
sns.countplot(x=df['issue_d'], order=df['issue_d'].value_counts().iloc[:7].index)
```

Out[42]:


```
<AxesSubplot:xlabel='issue_d', ylabel='count'>
```



As per this Dataset, Maximum loan has been issued on oct-2014

In [43]:

```
ax=sns.countplot(x=df['loan_status'])  
for label in ax.containers:  
    ax.bar_label(label)  
plt.show()
```



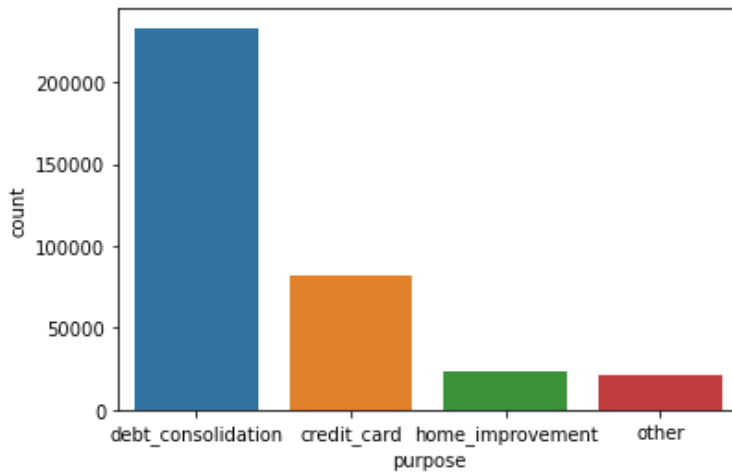
Around 80% of Loan has been Paid fully

In [44]:

```
sns.countplot(x=df['purpose'], order=df['purpose'].value_counts().iloc[:4].index)
```

Out[44]:

<AxesSubplot:xlabel='purpose', ylabel='count'>



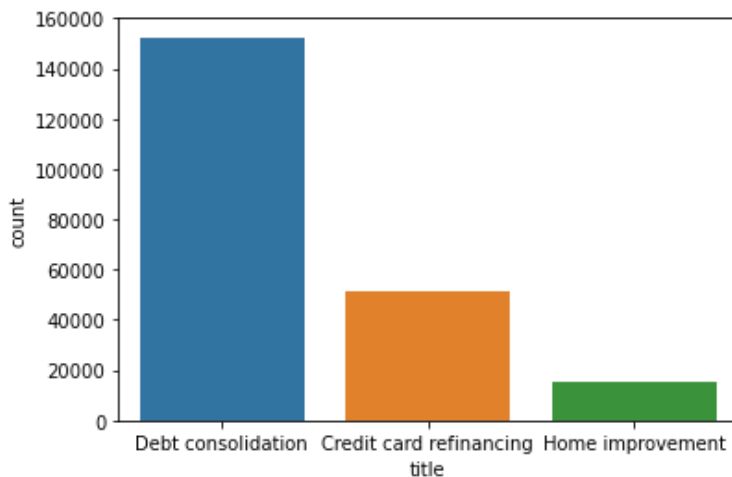
Debt consolidation and Credit card are the major purpose for loans

In [45]:

```
sns.countplot(x=df['title'], order=df['title'].value_counts().iloc[:3].index)
```

Out[45]:

<AxesSubplot:xlabel='title', ylabel='count'>




In [46]:

```
sns.boxplot(x=df['open_acc'])
```

Out[46]:

<AxesSubplot:xlabel='open_acc'>




0 20 40 60 80
open_acc

In [47]:

```
sns.countplot(x=df['initial_list_status'])
```

Out[47]:

<AxesSubplot:xlabel='initial_list_status', ylabel='count'>




In [48]:

```
sns.countplot(x=df['application_type'])
```

Out[48]:

<AxesSubplot:xlabel='application_type', ylabel='count'>




Almost all the loan is falling in Individual Category

In [49]:

```
sns.countplot(x=df['address'], order=df['address'].value_counts().iloc[:5].index)
```

Out[49]:

<AxesSubplot:xlabel='address', ylabel='count'>





Top 3 city of the borrower are AP,AE & AA


1.4.Bivariate Analysis

In [50]:

```
sns.boxplot(x=df['loan_amnt'], y=df['loan_status'])
```

Out[50]:

<AxesSubplot:xlabel='loan_amnt', ylabel='loan_status'>




Median Loan_amount of Fully_Paid loans is less tha Charged_off loans

In [51]:

```
sns.countplot(x=df['term (months)'], hue=df['loan_status'])
```

Out[51]:

<AxesSubplot:xlabel='term (months)', ylabel='count'>




In [52]:

```
sns.heatmap(data=pd.crosstab(df['term (months)'], df['loan_status'], normalize=True), annot=True)
```

Out[52]:

<AxesSubplot:xlabel='loan_status', ylabel='term (months)'>






The probability of full_loan_payment when term is 30 months is high as compared to 60 months

In [53]:

```
sns.boxplot(x=df['int_rate'], y=df['loan_status'])
```

Out[53]:


<AxesSubplot:xlabel='int_rate', ylabel='loan_status'>



Median Interest rate of fully_paid loan is less than Charged_off loans

In [54]:

```
g=sns.countplot(x=df['grade'], hue=df['loan_status'], order=['A', 'B', 'C', 'D', 'E', 'F', 'G'])
```





In [55]:

```
sns.heatmap(data=pd.crosstab(df['grade'], df['loan_status'], normalize=True), annot=True)
```

Out[55]:

<AxesSubplot:xlabel='loan_status', ylabel='grade'>






Probability of fully_paid for grade A is 93.75% followed grade B 87.54%

In [56]:

```
sns.countplot(x=df['emp_length (years)'], hue=df['loan_status'])
```

Out[56]:

<AxesSubplot:xlabel='emp_length (years)', ylabel='count'>




Ratio of fully_paid to charged_off is almost similar for different emp_lengths

In [57]:

```
ax=sns.boxplot(x=df['annual_inc'], y=df['loan_status'])
```

```
plt.xlim([0, 300000])
plt.show()
```



Median Annual income of the people who fully_paid loan is slightly high than people with Charged_off


In [58]:

```
sns.countplot(x=df['verification_status'], hue=df['loan_status'])
```

Out[58]:

<AxesSubplot:xlabel='verification_status', ylabel='count'>





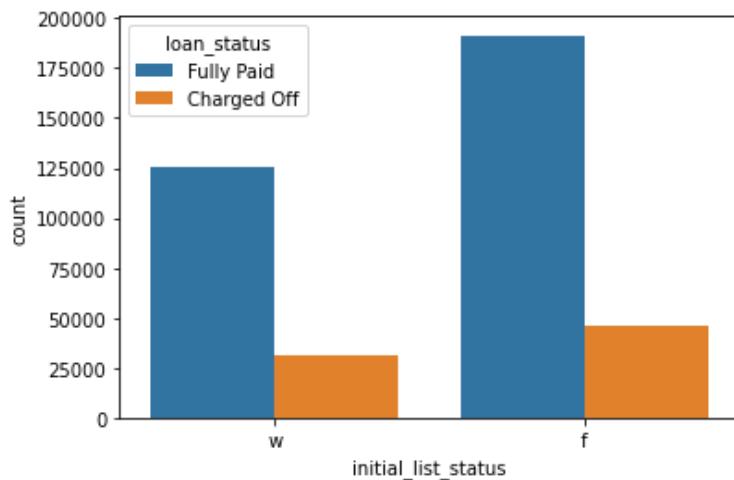
Loan Payment Status is independent with Verification status

In [59]:

```
sns.countplot(x=df['initial_list_status'], hue=df['loan_status'])
```

Out[59]:

<AxesSubplot:xlabel='initial_list_status', ylabel='count'>




Probability of full repayment is slightly high in case of 'f' tha 'w'

In [60]:

```
sns.regplot(x=df['loan_amnt'], y=df['installment'], fit_reg=True, scatter_kws={"color": "orange", 's': 1}, line_kws={"color": "blue"})
```

Out[60]:

<AxesSubplot:xlabel='loan_amnt', ylabel='installment'>




Installment is highly correlated with Loan_amount

In [61]:

```
df_emp_title=df.groupby('emp_title').mean('loan_amnt')
df_emp_title.reset_index(inplace=True)
df_emp_title=df_emp_title[['loan_amnt','emp_title']]
```


```
df_emp_title=df_emp_title.sort_values(by='loan_amnt',ascending=False)
df_emp_title.reset_index(drop=True,inplace=True)
sns.barplot(data=df_emp_title.iloc[:5,:],x='emp_title',y='loan_amnt')
plt.xticks(rotation=45)
plt.show()
```



cytotenetic tech, Logistic management analysts are having hogh mean loan_amount

In [62]:

```
df_address=df.groupby('address').mean('loan_amnt')
df_address.reset_index(inplace=True)
df_address=df_address[['loan_amnt','address']]
df_address=df_address.sort_values(by='loan_amnt',ascending=False)
df_address.reset_index(drop=True,inplace=True)
sns.barplot(data=df_address.iloc[:20,:],x='address',y='loan_amnt')
plt.xticks(rotation=45)
plt.show()
```



Not much difference in Mean_loan_amount with address


In [63]:

```
sns.scatterplot(x=df['loan_amnt'],y=df['installment'],hue=df['loan_status'])
```

Out[63]:

<AxesSubplot:xlabel='loan_amnt', ylabel='installment'>






Fully_Paid and Charged_off are having similar trend of installment v loan_amount

In [64]:

```
sns.boxplot(x=df['installment'], y=df['loan_status'])
```

Out[64]:

<AxesSubplot:xlabel='installment', ylabel='loan_status'>




Median of installment for Fully Paid is slightly less than Charged_off

In [65]:

```
sns.heatmap(data=df.corr())
```

Out[65]:

<AxesSubplot:>



Loan_amt is Highly correlated with installment annual income is correlated with installment and Loan_amt total_acc is correlated with open_acc pub_rec is correlated with pub_rec_bankruptcies

2.3.Outlier treatment

We have already observed in univariate analysis that we have outliers in loan amount, income and interest. We should be very mindful whenever dealing with outliers. There may be some higher or lower values in the data but at the same time, that could be the identity of the distribution. That's why I am not doing any treatment for outliers

2.4.Feature engineering

In [66]:

```
obj_cols = list(df.select_dtypes(include='object'))
obj_cols
```

Out[66]:

```
['term (months)',
 'grade',
 'sub_grade',
 'emp_title',
 'emp_length (years)',
 'home_ownership',
 'verification_status',
 'issue_d',
 'loan_status',
 'purpose',
 'title',
 'earliest_cr_line',
 'initial_list_status',
 'application_type',
 'address']
```

Converting 'term' and 'emp_length' to integer data type

In [67]:

```
df['term (months)']=df['term (months)'].astype(str).astype(int)
```

In [68]:

```
df['emp_length (years)']=df['emp_length (years)'].astype(str).astype(int)
```

Dropping date columns 'issue_d' & 'earliest_cr_line'

In [69]:

```
df.drop(columns=['issue_d'], inplace = True, axis=1)
```

In [70]:

```
df.drop(columns=['earliest_cr_line'], inplace = True, axis=1)
```

Dropping 'application_type' as it is having less variance(information) 99.9% are Individual

In [71]:

```
df.drop(columns=['application_type'], inplace = True, axis=1)
```

Converting Target('loan_status'), Ordinal Categorical to Categorical Numerical & the columns having many unique values ('title','emp_title','address','purpose')

In [72]:

```
col=['grade','sub_grade','loan_status','initial_list_status','address','title','emp_title',
    'address','purpose']
```

In [73]:


```
for i in col:
    df[i]=df[i].astype('category')
    df[i]=df[i].cat.codes
```

In [74]:

```
obj_col = list(df.select_dtypes(include='object'))
obj_col
```

Out[74]:

```
['home_ownership', 'verification_status']
```

One Hot Encoding for 'home_ownership','verification_status'

In [75]:

```
df=pd.concat((df, pd.get_dummies(df['home_ownership'])),axis=1)
df.drop(columns=['home_ownership'], inplace = True, axis=1)
```

In [76]:

```
df=pd.concat((df, pd.get_dummies(df['verification_status'])),axis=1)
df.drop(columns=['verification_status'], inplace = True, axis=1)
```

In [77]:

```
df.head()
```

Out[77]:

	loan_amnt	term (months)	int_rate	installment	grade	sub_grade	emp_title	emp_length (years)	annual_inc	loan_status	...	addr
70590	5000.0	36	15.61	174.83	3	15	65802	2	26000.0	1	...	
86600	12575.0	36	13.33	425.71	2	12	30284	10	127900.0	0	...	
169233	5125.0	36	16.78	182.16	2	14	4832	10	44316.0	1	...	
125243	8000.0	36	14.09	273.78	1	9	64835	4	105000.0	1	...	
315825	9600.0	36	10.99	314.25	1	7	54989	10	82940.0	1	...	

5 rows x 31 columns



In [78]:

```
df['loan_status'].value_counts()
```

Out[78]:

```
1    316271
0     77194
Name: loan_status, dtype: int64
```

1: Fully Paid, 0: Charged Off

2.5.Data preparation for modeling

In [79]:

```
X=df.drop(columns=['loan_status'])
```

In [80]:

```
X.head()
```

Out[80]:

	loan_amnt	term (months)	int_rate	installment	grade	sub_grade	emp_title	emp_length (years)	annual_inc	purpose	...	address
70590	5000.0	36	15.61	174.83	3	15	65802	2	26000.0	9	...	20
86600	12575.0	36	13.33	425.71	2	12	30284	10	127900.0	2	...	8
169233	5125.0	36	16.78	182.16	2	14	4832	10	44316.0	9	...	0
125243	8000.0	36	14.09	273.78	1	9	64835	4	105000.0	1	...	29
315825	9600.0	36	10.99	314.25	1	7	54989	10	82940.0	2	...	49

5 rows x 30 columns



In [81]:

```
Y=df[['loan_status']]
```

In [82]:

```
Y.head()
```

Out[82]:

	loan_status
70590	1
86600	0
169233	1
125243	1
315825	1

3.1.Build the Logistic Regression model and comment on the model statistics

In [83]:

```
from sklearn.model_selection import train_test_split

X_train_val, X_test, y_train_val, y_test = train_test_split(X,Y,test_size=0.2, random_state=42)

X_train, X_val, y_train, y_val = train_test_split(X_train_val,y_train_val,test_size=0.2, random_state=42)
```

In [84]:

```
print(X_train.shape, y_train.shape)
print(X_val.shape, y_val.shape)
print(X_test.shape, y_test.shape)

(251817, 30) (251817, 1)
(62955, 30) (62955, 1)
(78693, 30) (78693, 1)
```

In [85]:

```
from sklearn.preprocessing import StandardScaler

st = StandardScaler()

X_train_scaled = st.fit_transform(X_train.values)
X_val_scaled = st.transform(X_val.values)
X_test_scaled = st.transform(X_test.values)
```

In [86]:

```
y_train = y_train.values[:,0]
y_val = y_val.values[:,0]
y_test = y_test.values[:,0]
```

In [87]:

```
print(X_train_scaled.shape, y_train.shape)
print(X_val_scaled.shape, y_val.shape)
print(X_test_scaled.shape, y_test.shape)
```

```
(251817, 30) (251817,)
(62955, 30) (62955,)
(78693, 30) (78693,)
```

In [88]:

```
# import the class
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

acc=[]
for i in [0.001, 0.01, 0.1, 1, 10, 100, 1000]:
    # instantiate the model (using the default parameters)
    logreg = LogisticRegression(random_state=16,C=i)

    # fit the model with data
    logreg.fit(X_train_scaled, y_train)

    y_val_pred = logreg.predict(X_val_scaled)

    accuracy_score_val=accuracy_score(y_val, y_val_pred)
    acc.append(accuracy_score_val)
```

In [89]:

```
y_val_pred.shape
```

Out[89]:

```
(62955,)
```


In [90]:


```
plt.figure(figsize=(12, 6))
plt.plot(range(len(acc)), acc, color='red', linestyle='dashed', marker='o',
         markerfacecolor='blue', markersize=10)

plt.xlabel('Index of C-value List')
plt.ylabel('Accuracy Score')
```

Out[90]:

```
Text(0, 0.5, 'Accuracy Score')
```





AS maximum accuracy is at 1st index. Hence, We Can select HyperParameter C=0.001

In [99]:

```
logreg = LogisticRegression(random_state=16,C=0.001)

# fit the model with data
logreg.fit(X_train_scaled, y_train)

y_test_pred = logreg.predict(X_test_scaled)
```

3.2.Display model coefficients with column names

In [100]:

```
arr=logreg.coef_
arr
```

Out[100]:

```
array([[ -0.03509063,  -0.17477633,   0.04860077,  -0.04057853,  -0.13973434,
        -0.36622478,  -0.08597907,   0.02357343,   0.19511168,  -0.03241274,
         0.01403043,  -0.20539446,  -0.0995785 ,  -0.05081877,   0.04992771,
        -0.08436952,   0.08956595,  -0.00944975,   0.07265704,   0.0184871 ,
        -0.00113695,   0.01019199,   0.05683804,   0.00463742,  -0.00126526,
        -0.00430112,  -0.055442 ,   0.04843686,  -0.03513717,  -0.01246613]])
```

In [101]:

```
logreg.intercept_
```

Out[101]:

```
array([1.57525194])
```

In [102]:

```
abs_weight_ls=[]
col=list(X.columns)
print("Column : Weight")
for i in range(len(col)):
    w=str(round(arr[0][i],2))
    x=col[i]
    print(f"{x} : {w}")
    abs_weight_ls.append([x,abs(float(w))])
```

```
Column : Weight
loan_amnt : -0.04
term (months) : -0.17
int_rate : 0.05
installment : -0.04
grade : -0.14
sub_grade : -0.37
emp_title : -0.09
emp_length (years) : 0.02
annual_inc : 0.2
purpose : -0.03
title : 0.01
dti : -0.21
open_acc : -0.1
pub_rec : -0.05
revol_bal : 0.05
revol_util : -0.08
total_acc : 0.09
initial_list_status : -0.01
mort acc : 0.07
```

```
pub_rec_bankruptcies : 0.02
address : -0.0
ANY : 0.01
MORTGAGE : 0.06
NONE : 0.0
OTHER : -0.0
OWN : -0.0
RENT : -0.06
Not Verified : 0.05
Source Verified : -0.04
Verified : -0.01
```

In [103]:

```
abs_weight_ls.sort(key=lambda row:row[1],reverse=True)
```

In [107]:

```
abs_weight_ls[:10]
```

Out[107]:

```
[['sub_grade', 0.37],
 ['dti', 0.21],
 ['annual_inc', 0.2],
 ['term (months)', 0.17],
 ['grade', 0.14],
 ['open_acc', 0.1],
 ['emp_title', 0.09],
 ['total_acc', 0.09],
 ['revol_util', 0.08],
 ['mort_acc', 0.07]]
```

sub_grade, dti, annual_inc, term, grade, open_acc, total_acc, emp_title are most contributing Feature

4.1.ROC AUC Curve & comments

In [105]:

```
from sklearn.metrics import plot_roc_curve
```

In [106]:


```
plot_roc_curve(logreg,X_test_scaled,y_test)
```

C:\Users\rahul.kumar\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

```
warnings.warn(msg, category=FutureWarning)
```

Out[106]:

```
<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x267303f9a90>
```



ROC AUC Score is 0.71, the reason for low roc-auc is imbalance data


4.2.Precision Recall Curve & comments

In [108]:

```
from sklearn.metrics import precision_recall_curve
```

In [114]:

```
precision, recall, thresholds=precision_recall_curve(y_test,y_test_pred)
plt.fill_between(recall, precision)
plt.ylabel("Precision")
plt.xlabel("Recall")
plt.title("Train Precision-Recall curve");
```



we are getting 0.9 as Precision-Recall Value, which is suggesting good model

4.3.Classification Report (Confusion Matrix etc)

In [116]:

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

In [117]:

```
# Calculate the accuracy
accuracy = accuracy_score(y_test, y_test_pred)

# Calculate the precision
precision = precision_score(y_test, y_test_pred)

# Calculate the recall
recall = recall_score(y_test, y_test_pred)

# Calculate the f1 score
f1 = f1_score(y_test, y_test_pred)

# Print the results
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
Accuracy: 0.8065012135768111
Precision: 0.8139659545537745
Recall: 0.9842225910995178
F1 Score: 0.8910341274214439
```

In [120]:


```
from sklearn.metrics import confusion_matrix
```

In [124]:

```
# Create a confusion matrix
sns.heatmap(confusion_matrix(y_test, y_test_pred),annot=True, fmt=".1f")
```

Out[124]:

<AxesSubplot:>



4.4.1.How can we make sure that our model can detect real defaulters and there are less false positives? This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.

As our recall for class 0 is 90%, we can detect it, if we want further improvement then we need to change the value of threshold of probability > 0.5 (for true class)

4.4.2.Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone.

No, we should give loans but with minimal risk and not with 0 risk. If we take 0 risk, then no point of doing business in loan sector

4.5.Actionable Insights & Recommendations

1. sub_grade, dti, annual_inc, term, grade, open_acc, total_acc, emp_title are most contributing Feature. 2. Address is not contributing to the outcome. If we remove from feature, we will get slightly better performance time. 3. As the data is imbalance, Upscaling and Downscaling may lead to more accurate results. 4. Other Models like KNN or Decision Tree may improve the accuracy or F1. 5. We need to improve Recall for Class 0(not fully paid) to drive business with minimal risk.1.What percentage of customers have fully paid their Loan Amount? 80% 2.Comment about the correlation between Loan Amount and Installment features? Highly Co-relarelated 3.The majority of people have home ownership as _____? Mortgage 4.People with grades 'A' are more likely to fully pay their loan (T/F)? True 5.Name the top 2 afforded job titles? Teacher & Manager 6.Thinking from a bank's perspective, which metric should our primary focus be on ROC AUC Precision Recall F1 Score? F1 Score 7.How does the gap in precision and recall affect the bank? Recall could be the most important criteria, spetially for class 0. 8.Which were the features that heavily affected the outcome? sub_grade,dti,annual_inc,term, grade,open_acc,total_acc, emp_title 9.Will the results be affected by geographical location? (Yes/No) No